12-Month Business Case — Family Case Review (Sep 2026 – Aug 2027)

Status: For decision · **Owner:** Product/Finance · **Prepared:** 2026-08-30 **Sources:** `snotnoselegal_market_study_mvp_gtm` (market frame §1, pricing §6, unit economics §7, GTM §8) · `mvp_v1_prd.md` · `product_roadmap.md` · `../implementation/mvp_v1_implementation_plan.md` (build status & timeline) · PO decisions of Aug 2026 (SLA, pricing, brand, tax, stack) · **live-measured pipeline costs from the first real reference-case runs (Aug 30, 2026)**

---

1. Executive summary

Snot Nose Legal's **Family Case Review** sells a $299 flat-fee, AI-analyzed, human-QA'd review of a Texas criminal court record to the one persona in the post-conviction market no incumbent serves and who actually pays: the inmate's family. The competitive anchor is not software but the **~$3,000 attorney file review** and the $0 pro se path that burns a family's only writ shot. The market study establishes genuine white space (habeas-specific analysis; consumer access) that none of the five incumbent platforms can enter without abandoning their per-professional-seat licensing.

This case models the next 12 months: four months to launch-ready (per the implementation plan; **the majority of the engineering is already built and integration-tested** — auth, payments, digitization on AWS Textract, the Claude Opus 5 analysis pipeline, QA console, and consumer surfaces are live in the repo as of this writing) and eight months of operation.

The headline numbers (base case): first-year revenue ≈ **$36k**, peak cumulative cash need ≈ **$78k**, monthly contribution turns positive around **~30 cases/month** at current cost structure — i.e., the business does not self-fund inside 12 months at bootstrap staffing; it buys three things instead: a **proven engine on paid real records**, the **reference metrics** (recall, refund rate, CAC by channel) that de-risk every later tier, and the **referral flywheel** into the Justice/Advocate tiers where the durable revenue lives (roadmap v2/v3). The recommendation (§10) is to proceed, with the pre-agreed kill/pivot checkpoints in §8.

2. Problem & product

When a Texas felony conviction becomes final, the family faces a brutal information gap: pay ~$3,000 for an attorney merely to read the file, pay $15k–$50k+ for representation on faith, or let the inmate file pro se — where a weak application burns the effective one shot Texas's subsequent-writ bar allows. **The product sells the decision, not the writ**: every finding grounded to page and line, screened for the five claim families (preserved error, IAC, Brady, 11.073 junk science, sentencing/deadlines), passed through a human QA gate, and delivered as Part A (plain English) + Part B (attorney-ready packet).

Trust is the product. The engineering enforces it structurally: findings whose quotes aren't verbatim in the record are dropped by a hard filter; citations are hash-anchored and re-verified at every render; nothing reaches a customer unreviewed; deletion, refunds, and disclosures are audited operations, not promises.

3. Market

Measure	Size	Definition
TAM	$30B	Global legal tech & legal AI
SAM	$1.5–2B	US criminal defense / public defender legal tech
SOM	$50–100M	Texas post-conviction — the nation's highest Art. 11.07 writ volume

The consumer wedge is deliberately unsized by incumbents: all five platforms license to lawyers. Texas structural facts driving demand: IAC is the dominant claim; most 11.07 applications are filed pro se; the AEDPA one-year clock is miscalculated fatally and computed by no platform; decades-old scans break incumbents' tooling (our Textract pipeline with per-page confidence is built for exactly these — proven this week on a real 602-page reporter's record).

Competitive floor: research is free to the Texas bar (vLex), document review is free to nonprofits (Everlaw for Good). The product never positions as either.

4. Pricing & offer (as decided, Aug 2026)

Offer	Price	Notes
Family Case Review	**$299 one-time**	Uniform across trial and plea lanes (PO decision); 5,000-page cap
Page overage	+$49 / 2,500 pages	Charged in-flow only when the count crosses the line
Re-run with new documents	$99	After report delivery
Attorney-signed review add-on	~$499	v1.1 (from ~May 2027); licensed TX attorney per review
Payment options	Stripe-native installments (Affirm/Klarna)	Stripe Tax enabled; TX data-processing determination pending

10-business-day SLA from records-complete. Later tiers (Advocate $99/seat, Chambers $249/seat, Justice $0-to-cap, Sovereign site license) are out of this case's revenue but are its strategic payoff.

5. Unit economics

Study budget vs. what the pipeline **actually measured** on the reference case this week:

Cost line	Study budget	Measured (Aug 30, 2026)	Note
OCR	~$5 / 3,000 pp	**~$0.90** (602 pp × Textract $1.50/1k)	scales ~$7.50 at the 5,000-page cap
LLM analysis	~$12	**~$3.30/run** (Claude Opus 5, 4 screens, 264K-token record; prompt cache verified at 0.1× reads)	headroom for the 5-screen set + adjudication
Embeddings + storage	~$1	S3 pennies; embeddings paused by design	
Payment processing	~$9	2.9% + $0.30	+ 0.5% Stripe Tax when taxable
Human QA	~$12.50	30 min × $25/hr loaded	every report
Refund/chargeback reserve	~$15	5% of revenue	consumer stress purchases
Total COGS	**~$54**	**tracking ≤ $45 at current scale**	**≥ 82% gross margin at $299**

This case models the conservative **$54**; the measured trend is upside.

6. 12-month plan

Phase	Months	Contents	Source
Finish build → launch gates	Sep–Dec 2026	Remaining M4–M7 work (batch/adjudication, deadline-engine vectors, PDF render, analytics); **eval harness green on both reference cases; attorney review; E&O bound**	implementation plan §3/§7
Launch	Jan 2027	TX-only, SEO + prison-family communities + clinic-declined referrals first; paid search as a capped test	market study §8.2
Ramp	Feb–Apr 2027	90-day target ~10 paid cases/mo; S0 outcome-mix report validates market sizing	analytics plan §5
v1.1	~Apr–May 2027	Spanish Part A + purchase flow; attorney-signed add-on; A/V add-on priced	roadmap v1.1
Justice-tier seeding	Jun–Aug 2027	Consented high-viability referrals to 2–3 clinics — the v2 gate	roadmap v2

7. Financial model — three scenarios

Assumptions (labeled): bootstrap staffing (founder unpaid + part-time senior contractor: $8k/mo build, $6k/mo operating — the funded-team alternative per the implementation plan runs ~$30k/mo and is shown in §8); fixed infra + E&O ~$600/mo (study §7.4); counsel $4k one-time (Nov, gate 2) + $500/mo retainer from launch; marketing = max($500/mo test budget, blended CAC $60 on the ~50% of cases from paid channels — organic channels are the durable base); overage attach 10%, re-run 5%; attorney-signed add-on 8% attach from May 2027 (base/upside only) at $250 attorney cost; volumes per the study's break-even ladder with the 90-day ~10/mo target hit in March (base).

Conservative scenario

Month	Cases	Revenue	COGS	Gross	Opex	EBITDA	Cumulative
Sep-26	0	$0	$0	$0	$8,600	$-8,600	$-8,600
Oct-26	0	$0	$0	$0	$8,600	$-8,600	$-17,200
Nov-26	0	$0	$0	$0	$12,600	$-12,600	$-29,800
Dec-26	0	$0	$0	$0	$8,600	$-8,600	$-38,400
Jan-27	2	$598	$108	$490	$7,600	$-7,110	$-45,510
Feb-27	3	$897	$162	$735	$7,600	$-6,865	$-52,375
Mar-27	5	$1,495	$270	$1,225	$7,600	$-6,375	$-58,750
Apr-27	7	$2,142	$378	$1,764	$7,600	$-5,836	$-64,586
May-27	9	$2,740	$486	$2,254	$7,600	$-5,346	$-69,932
Jun-27	11	$3,437	$594	$2,843	$7,600	$-4,757	$-74,689
Jul-27	13	$4,035	$702	$3,333	$7,600	$-4,267	$-78,956
Aug-27	15	$4,682	$810	$3,872	$7,600	$-3,728	$-82,684

Year totals — revenue $20,026, EBITDA $-82,684; peak cumulative cash need $82,684; first EBITDA-positive month: not in window.

Base scenario

Month	Cases	Revenue	COGS	Gross	Opex	EBITDA	Cumulative
Sep-26	0	$0	$0	$0	$8,600	$-8,600	$-8,600
Oct-26	0	$0	$0	$0	$8,600	$-8,600	$-17,200
Nov-26	0	$0	$0	$0	$12,600	$-12,600	$-29,800
Dec-26	0	$0	$0	$0	$8,600	$-8,600	$-38,400
Jan-27	3	$897	$162	$735	$7,600	$-6,865	$-45,265
Feb-27	6	$1,843	$324	$1,519	$7,600	$-6,081	$-51,346
Mar-27	10	$3,039	$540	$2,499	$7,600	$-5,101	$-56,447
Apr-27	13	$4,035	$702	$3,333	$7,600	$-4,267	$-60,714
May-27	17	$5,779	$1,168	$4,611	$7,610	$-2,999	$-63,713
Jun-27	21	$7,474	$1,634	$5,840	$7,730	$-1,890	$-65,603
Jul-27	25	$8,670	$1,850	$6,820	$7,850	$-1,030	$-66,633
Aug-27	30	$10,313	$2,120	$8,193	$8,000	$193	$-66,440

Year totals — revenue $42,050, EBITDA $-66,440; peak cumulative cash need $66,633; first EBITDA-positive month: Aug-27.

Upside scenario

Month	Cases	Revenue	COGS	Gross	Opex	EBITDA	Cumulative
Sep-26	0	$0	$0	$0	$8,600	$-8,600	$-8,600
Oct-26	0	$0	$0	$0	$8,600	$-8,600	$-17,200
Nov-26	0	$0	$0	$0	$12,600	$-12,600	$-29,800
Dec-26	0	$0	$0	$0	$8,600	$-8,600	$-38,400
Jan-27	5	$1,495	$270	$1,225	$7,600	$-6,375	$-44,775
Feb-27	9	$2,740	$486	$2,254	$7,600	$-5,346	$-50,121
Mar-27	14	$4,334	$756	$3,578	$7,600	$-4,022	$-54,143
Apr-27	19	$5,878	$1,026	$4,852	$7,670	$-2,818	$-56,961
May-27	25	$8,670	$1,850	$6,820	$7,850	$-1,030	$-57,991
Jun-27	31	$10,612	$2,174	$8,438	$8,030	$408	$-57,583
Jul-27	38	$13,253	$2,802	$10,451	$8,240	$2,211	$-55,372
Aug-27	45	$15,845	$3,430	$12,415	$8,450	$3,965	$-51,407

Year totals — revenue $62,827, EBITDA $-51,407; peak cumulative cash need $57,991; first EBITDA-positive month: Jun-27.

8. Cash need, staffing sensitivity, and checkpoints

- **Funding requirement (bootstrap):** peak cumulative need ≈ **$78k (base)** / $83k (conservative) / $72k (upside). Round to a **$100k envelope** for working-capital comfort (chargebacks, Textract/LLM spikes, counsel overruns).
- **Funded-team alternative:** at the implementation plan's full staffing (~$30k/mo), the same 12 months need ≈ **$320–350k**. The bootstrap path is viable precisely because the build is largely done; choose funded staffing only if v1.1/v2 acceleration is worth buying.
- **Sensitivity (base case):** CAC doubling to $120 costs ~$5k/yr (manageable — contribution/case ≈ $245); COGS at the measured ~$45 adds ~$1.5k/yr; a $349 price test (+17%) adds ~$8k/yr revenue at unknown conversion cost — run it only through the §4.3-compliant experiment path. The model is **volume-dominated**: the single number that matters is paid cases/month.
- **Checkpoints (pre-agreed):**
- **Dec 2026:** eval harness green + attorney sign-off, or launch slips — do not launch on an unproven engine.
- **Mar 2027 (90 days post-launch):** ≥ 8–10 cases/mo and refund rate ≤ 5% → continue; materially below → the S0 outcome-mix data decides pivot (plea-lane emphasis, price, or channel) before more spend.
- **Jun 2027:** ≥ 20% referral opt-in → begin Justice-tier seeding (the v2 gate); below → double down on B2C channels first.

9. Risks (top 5, from SWOT + study §8.4 + risk register)

Risk	Mitigation
UPL exposure on the consumer tier	Information-report framing throughout (built); counsel gate 2; attorney-signed add-on keeps a human path; never route toward pro se filing
Volume doesn't materialize (the model's dominant risk)	Channels are cheap and compounding (SEO, communities, clinic referrals); checkpoints cap the downside; contribution ≈ $245/case means even small volume is not cash-destructive
Chargebacks/refunds above 5%	W-2 disclosure archive per case (built — the E-6 evidence packet is one query); 10-day SLA honesty; reserve modeled
Incumbent module entry	The moat is the white-space spine + consumer trust surface; incumbents' seat-licensing makes B2C structurally awkward for them
Model-quality regression / eval not green	Recall-first harness on attorney-labeled ledgers is a launch **gate**, not a hope; today's live runs already shook out real pipeline defects pre-launch

10. KPIs (weekly, from the analytics plan)

Landing → check ≥ 25% · check completion ≥ 80% · fit → ack ≥ 60% · ack → paid ≥ 70% · **S0 outcome mix** (the market-validation panel) · time-to-records-complete · QA rejection < 10% · refund ≤ 5% · COGS/case ≤ $54 · north star: **completed reviews that reach a lawyer**.

11. Recommendation

Proceed on the bootstrap path with a $100k working-capital envelope. The first year is an investment year by design: the P&L loss (~$78k base) purchases a proven, eval-gated analysis engine on paid real records, channel economics measured against the study's targets, and the pre-triaged referral pipeline that gates the Justice tier — the sequence the study identifies as the only durable route to the $50–100M Texas SOM. The two disciplines that keep the downside bounded: the launch gates are hard (no eval green, no launch), and the March volume checkpoint decides pivot with data, not hope.