

12-Month Business Case — Family Case Review (Sep 2026 – Aug 2027)

Status: For decision · **Owner:** Product/Finance · **Prepared:** 2026-08-30 **Sources:**

snotnoselegal_market_study_mvp_gtm (market frame \$1, pricing \$6, unit economics \$7, GTM \$8) · mvp_v1_prd.md · product_roadmap.md · ../implementation/mvp_v1_implementation_plan.md (build status & timeline) · PO decisions of Aug 2026 (SLA, pricing, brand, tax, stack) · **live-measured pipeline costs from the first real reference-case runs (Aug 30, 2026)**

1. Executive summary

Snot Nose Legal's **Family Case Review** sells a \$299 flat-fee, AI-analyzed, human-QA'd review of a Texas criminal court record to the one persona in the post-conviction market no incumbent serves and who actually pays: the inmate's family. The competitive anchor is not software but the **~\$3,000 attorney file review** and the \$0 pro se path that burns a family's only writ shot. The market study establishes genuine white space (habeas-specific analysis; consumer access) that none of the five incumbent platforms can enter without abandoning their per-professional-seat licensing.

This case models the next 12 months: four months to launch-ready (per the implementation plan; **the majority of the engineering is already built and integration-tested** — auth, payments, digitization on AWS Textract, the Claude Opus 5 analysis pipeline, QA console, and consumer surfaces are live in the repo as of this writing) and eight months of operation.

The headline numbers (base case): first-year revenue ≈ **\$36k**, peak cumulative cash need ≈ **\$78k**, monthly contribution turns positive around **~30 cases/month** at current cost structure — i.e., the business does not self-fund inside 12 months at bootstrap staffing; it buys three things instead: a **proven engine on paid real records**, the **reference metrics** (recall, refund rate, CAC by channel) that de-risk every later tier, and the **referral flywheel** into the Justice/Advocate tiers where the durable revenue lives (roadmap v2/v3). The recommendation (\$10) is to proceed, with the pre-agreed kill/pivot checkpoints in §8.

2. Problem & product

When a Texas felony conviction becomes final, the family faces a brutal information gap: pay ~\$3,000 for an attorney merely to read the file, pay \$15k–\$50k+ for representation on faith, or let the inmate file pro se — where a weak application burns the effective one shot Texas's subsequent-writ bar allows. **The product sells the decision, not the writ:** every finding grounded to page and line, screened for the five claim families (preserved error, IAC, Brady, 11.073 junk science, sentencing/deadlines), passed through a human QA gate, and delivered as Part A (plain English) + Part B (attorney-ready packet).

Trust is the product. The engineering enforces it structurally: findings whose quotes aren't verbatim in the record are dropped by a hard filter; citations are hash-anchored and re-verified at every render; nothing reaches a customer unreviewed; deletion, refunds, and disclosures are audited operations, not promises.

3. Market

Measure	Size	Definition
TAM	\$30B	Global legal tech & legal AI
SAM	\$1.5–2B	US criminal defense / public defender legal tech
SOM	\$50–100M	Texas post-conviction — the nation's highest Art. 11.07 writ volume

The consumer wedge is deliberately unsized by incumbents: all five platforms license to lawyers. Texas structural facts driving demand: IAC is the dominant claim; most 11.07 applications are filed pro se; the AEDPA one-year clock is miscalculated fatally and computed by no platform; decades-old scans break incumbents' tooling (our Textract pipeline with per-page confidence is built for exactly these — proven this week on a real 602-page reporter's record).

Competitive floor: research is free to the Texas bar (vLex), document review is free to nonprofits (Everlaw for Good). The product never positions as either.

4. Pricing & offer (as decided, Aug 2026)

Offer	Price	Notes
Family Case Review	\$299 one-time	Uniform across trial and plea lanes (PO decision); 5,000-page cap
Page overage	+\$49 / 2,500 pages	Charged in-flow only when the count crosses the line
Re-run with new documents	\$99	After report delivery
Attorney-signed review add-on	~\$499	v1.1 (from ~May 2027); licensed TX attorney per review
Payment options	Stripe-native installments (Affirm/Klarna)	Stripe Tax enabled; TX data-processing determination pending

10-business-day SLA from records-complete. Later tiers (Advocate \$99/seat, Chambers \$249/seat, Justice \$0-to-cap, Sovereign site license) are out of this case's revenue but are its strategic payoff.

5. Unit economics

Study budget vs. what the pipeline **actually measured** on the reference case this week:

Cost line	Study budget	Measured (Aug 30, 2026)	Note
OCR	~\$5 / 3,000 pp	~ \$0.90 (602 pp × Textract \$1.50/1k)	scales ~\$7.50 at the 5,000-page cap
LLM analysis	~\$12	~ \$3.30/run (Claude Opus 5, 4 screens, 264K-token record; prompt cache verified at 0.1× reads)	headroom for the 5-screen set + adjudication
Embeddings + storage	~\$1	S3 pennies; embeddings paused by design	
Payment processing	~\$9	2.9% + \$0.30	+ 0.5% Stripe Tax when taxable
Human QA	~\$12.50	30 min × \$25/hr loaded	every report
Refund/chargeback reserve	~\$15	5% of revenue	consumer stress purchases
Total COGS	~\$54	tracking ≤ \$45 at current scale	≥ 82% gross margin at \$299

This case models the conservative **\$54**; the measured trend is upside.

6. 12-month plan

Phase	Months	Contents	Source
Finish build → launch gates	Sep–Dec 2026	Remaining M4–M7 work (batch/adjudication, deadline-engine vectors, PDF render, analytics); eval harness green on both reference cases; attorney review; E&O bound	implementation plan \$3/\$7
Launch	Jan 2027	TX-only, SEO + prison-family communities + clinic-declined referrals first; paid search as a capped test	market study \$8.2
Ramp	Feb–Apr 2027	90-day target ~10 paid cases/mo; S0 outcome-mix report validates market sizing	analytics plan \$5
v1.1	~Apr–May 2027	Spanish Part A + purchase flow; attorney-signed add-on; A/V add-on priced	roadmap v1.1
Justice-tier seeding	Jun–Aug 2027	Consented high-viability referrals to 2–3 clinics — the v2 gate	roadmap v2

7. Financial model — three scenarios

Assumptions (labeled): bootstrap staffing (founder unpaid + part-time senior contractor: \$8k/mo build, \$6k/mo operating — the funded-team alternative per the implementation plan runs ~\$30k/mo and is shown in \$8); fixed infra + E&O ~\$600/mo (study \$7.4); counsel \$4k one-time (Nov, gate 2) + \$500/mo retainer from launch; marketing = max(\$500/mo test budget, blended CAC \$60 on the ~50% of cases from paid channels — organic channels are the durable base); overage attach 10%, re-run 5%; attorney-signed add-on 8% attach from May 2027 (base/upside only) at \$250 attorney cost; volumes per the study's break-even ladder with the 90-day ~10/mo target hit in March (base).

Conservative scenario

Month	Cases	Revenue	COGS	Gross	Opex	EBITDA	Cumulative
Sep-26	0	\$0	\$0	\$0	\$8,600	\$-8,600	\$-8,600
Oct-26	0	\$0	\$0	\$0	\$8,600	\$-8,600	\$-17,200
Nov-26	0	\$0	\$0	\$0	\$12,600	\$-12,600	\$-29,800
Dec-26	0	\$0	\$0	\$0	\$8,600	\$-8,600	\$-38,400
Jan-27	2	\$598	\$108	\$490	\$7,600	\$-7,110	\$-45,510
Feb-27	3	\$897	\$162	\$735	\$7,600	\$-6,865	\$-52,375
Mar-27	5	\$1,495	\$270	\$1,225	\$7,600	\$-6,375	\$-58,750
Apr-27	7	\$2,142	\$378	\$1,764	\$7,600	\$-5,836	\$-64,586
May-27	9	\$2,740	\$486	\$2,254	\$7,600	\$-5,346	\$-69,932
Jun-27	11	\$3,437	\$594	\$2,843	\$7,600	\$-4,757	\$-74,689
Jul-27	13	\$4,035	\$702	\$3,333	\$7,600	\$-4,267	\$-78,956
Aug-27	15	\$4,682	\$810	\$3,872	\$7,600	\$-3,728	\$-82,684

Year totals — revenue \$20,026, EBITDA \$-82,684; peak cumulative cash need \$82,684; first EBITDA-positive month: not in window.

Base scenario

Month	Cases	Revenue	COGS	Gross	Opex	EBITDA	Cumulative
Sep-26	0	\$0	\$0	\$0	\$8,600	\$-8,600	\$-8,600
Oct-26	0	\$0	\$0	\$0	\$8,600	\$-8,600	\$-17,200
Nov-26	0	\$0	\$0	\$0	\$12,600	\$-12,600	\$-29,800
Dec-26	0	\$0	\$0	\$0	\$8,600	\$-8,600	\$-38,400
Jan-27	3	\$897	\$162	\$735	\$7,600	\$-6,865	\$-45,265
Feb-27	6	\$1,843	\$324	\$1,519	\$7,600	\$-6,081	\$-51,346
Mar-27	10	\$3,039	\$540	\$2,499	\$7,600	\$-5,101	\$-56,447
Apr-27	13	\$4,035	\$702	\$3,333	\$7,600	\$-4,267	\$-60,714
May-27	17	\$5,779	\$1,168	\$4,611	\$7,610	\$-2,999	\$-63,713
Jun-27	21	\$7,474	\$1,634	\$5,840	\$7,730	\$-1,890	\$-65,603
Jul-27	25	\$8,670	\$1,850	\$6,820	\$7,850	\$-1,030	\$-66,633
Aug-27	30	\$10,313	\$2,120	\$8,193	\$8,000	\$193	\$-66,440

Year totals — revenue \$42,050, EBITDA \$-66,440; peak cumulative cash need \$66,633; first EBITDA-positive month: Aug-27.

Upside scenario

Month	Cases	Revenue	COGS	Gross	Opex	EBITDA	Cumulative
Sep-26	0	\$0	\$0	\$0	\$8,600	\$-8,600	\$-8,600
Oct-26	0	\$0	\$0	\$0	\$8,600	\$-8,600	\$-17,200
Nov-26	0	\$0	\$0	\$0	\$12,600	\$-12,600	\$-29,800
Dec-26	0	\$0	\$0	\$0	\$8,600	\$-8,600	\$-38,400
Jan-27	5	\$1,495	\$270	\$1,225	\$7,600	\$-6,375	\$-44,775
Feb-27	9	\$2,740	\$486	\$2,254	\$7,600	\$-5,346	\$-50,121
Mar-27	14	\$4,334	\$756	\$3,578	\$7,600	\$-4,022	\$-54,143
Apr-27	19	\$5,878	\$1,026	\$4,852	\$7,670	\$-2,818	\$-56,961
May-27	25	\$8,670	\$1,850	\$6,820	\$7,850	\$-1,030	\$-57,991
Jun-27	31	\$10,612	\$2,174	\$8,438	\$8,030	\$408	\$-57,583
Jul-27	38	\$13,253	\$2,802	\$10,451	\$8,240	\$2,211	\$-55,372
Aug-27	45	\$15,845	\$3,430	\$12,415	\$8,450	\$3,965	\$-51,407

Year totals — revenue \$62,827, EBITDA \$-51,407; peak cumulative cash need \$57,991; first EBITDA-positive month: Jun-27.

8. Cash need, staffing sensitivity, and checkpoints

- **Funding requirement (bootstrap):** peak cumulative need $\approx$ **\$78k (base)** / \$83k (conservative) / \$72k (upside). Round to a **\$100k envelope** for working-capital comfort (chargebacks, Textract/LLM spikes, counsel overruns).
- **Funded-team alternative:** at the implementation plan's full staffing ($\sim$ \$30k/mo), the same 12 months need $\approx$ **\$320–350k**. The bootstrap path is viable precisely because the build is largely done; choose funded staffing only if v1.1/v2 acceleration is worth buying.
- **Sensitivity (base case):** CAC doubling to \$120 costs $\sim$ \$5k/yr (manageable — contribution/case $\approx$ \$245); COGS at the measured $\sim$ \$45 adds $\sim$ \$1.5k/yr; a \$349 price test (+17%) adds $\sim$ \$8k/yr revenue at unknown conversion cost — run it only through the \$4.3-compliant experiment path. The model is **volume-dominated**: the single number that matters is paid cases/month.
- **Checkpoints (pre-agreed):**
- **Dec 2026:** eval harness green + attorney sign-off, or launch slips — do not launch on an unproven engine.
- **Mar 2027 (90 days post-launch):** $\geq$ 8–10 cases/mo and refund rate $\leq$ 5% $\rightarrow$ continue; materially below $\rightarrow$ the S0 outcome-mix data decides pivot (plea-lane emphasis, price, or channel) before more spend.
- **Jun 2027:** $\geq$ 20% referral opt-in $\rightarrow$ begin Justice-tier seeding (the v2 gate); below $\rightarrow$ double down on B2C channels first.

8b. The path to profitability (BA answer to “how do we make money on this”)

The base case loses money for exactly one structural reason: **fixed operating cost (\$7.6k/mo) needs ~31 cases/month to cover at \$245 contribution, and the ramp only reaches 30 by month 12.** Profitability is therefore engineered from three levers, in order of leverage:

Lever 1 — lean operations (the big one). The \$6k/mo operating contractor is 79% of post-launch opex, budgeted for a build that is *already done*. Founder-run operations with an on-call contractor (~\$1.5k/mo) cut fixed cost to ~\$3.1k/mo → **break-even falls from 31 to 13 cases/month** — inside the study’s own 90-day trajectory.

Lever 2 — contribution stacking (no new customers needed). Measured COGS (~\$45, not \$54) + overage/re-run attach + the v1.1 attorney-signed add-on pulled forward (10% attach × \$249 margin) lifts contribution from \$245 to **~\$289/case** → break-even **~11 cases/month**.

Lever 3 — Advocate-tier early access (recurring revenue, warm leads). The QA console is already the attorney workspace dogfood, and the referral list produces warm attorney contacts from month one. An early-access Advocate seat at \$99/mo needs no new engineering to pilot: **10 seats ≈ 3 cases of monthly contribution, 25 seats ≈ 9 — recurring.** This is the roadmap’s Phase-3 revenue sampled early, without gating v2.

Re-modeled (lean ops + stacked contribution, same base volumes):

Post-launch month	Cases	EBITDA	Cumulative
Jan-27	3	−\$2,338	−\$40,738
Feb-27	6	−\$1,527	−\$42,265
Mar-27	10	−\$511	−\$42,776
Apr-27	13	+\$350	−\$42,426
May-27	17	+\$1,903	−\$40,523
Jun-27	21	+\$2,799	−\$37,724
Jul-27	25	+\$3,695	−\$34,029
Aug-27	30	+\$5,212	−\$28,817

Monthly EBITDA turns positive in April 2027 (month 4 post-launch), peak cash need drops from ~\$78k to **~\$43k**, and the exit run-rate (~\$5k/mo EBITDA and climbing, before any Advocate seats) pays back the build burn during year two. Add 25 early-access Advocate seats and break-even arrives a month earlier with ~\$2.5k/mo of recurring on top.

What this does NOT change: the volume dependency. Every lever above lowers the bar; only the channels (clinic-declined referrals at zero CAC, SEO, communities) clear it. The March checkpoint in §8 remains the decision that matters.

9. Risks (top 5, from SWOT + study §8.4 + risk register)

Risk	Mitigation
UPL exposure on the consumer tier	Information-report framing throughout (built); counsel gate 2; attorney-signed add-on keeps a human path; never route toward pro se filing
Volume doesn't materialize (the model's dominant risk)	Channels are cheap and compounding (SEO, communities, clinic referrals); checkpoints cap the downside; contribution ≈ \$245/case means even small volume is not cash-destructive
Chargebacks/refunds above 5%	W-2 disclosure archive per case (built — the E-6 evidence packet is one query); 10-day SLA honesty; reserve modeled
Incumbent module entry	The moat is the white-space spine + consumer trust surface; incumbents' seat-licensing makes B2C structurally awkward for them
Model-quality regression / eval not green	Recall-first harness on attorney-labeled ledgers is a launch gate , not a hope; today's live runs already shook out real pipeline defects pre-launch

10. KPIs (weekly, from the analytics plan)

Landing → check ≥ 25% · check completion ≥ 80% · fit → ack ≥ 60% · ack → paid ≥ 70% · **S0 outcome mix** (the market-validation panel) · time-to-records-complete · QA rejection < 10% · refund ≤ 5% · COGS/case ≤ \$54 · north star: **completed reviews that reach a lawyer**.

11. Recommendation

Proceed on the bootstrap path with a \$100k working-capital envelope. The first year is an investment year by design: the P&L loss (~\$78k base) purchases a proven, eval-gated analysis engine on paid real records, channel economics measured against the study's targets, and the pre-triaged referral pipeline that gates the Justice tier — the sequence the study identifies as the only durable route to the \$50–100M Texas SOM. The two disciplines that keep the downside bounded: the launch gates are hard (no eval green, no launch), and the March volume checkpoint decides pivot with data, not hope.