

# 12-Month Business Case — Family Case Review (Sep 2026 – Aug 2027)

---

**Status:** For decision · **Owner:** Product/Finance · **Prepared:** 2026-08-30 · **Updated:** 2026-09-01 **Sources:** snotnoselegal\_market\_study\_mvp\_gtm (market frame \$1, pricing \$6, unit economics \$7, GTM \$8) · mvp\_v1\_prd.md · product\_roadmap.md · ../implementation/mvp\_v1\_implementation\_plan.md (build status & timeline) · PO decisions of Aug 2026 (SLA, pricing, brand, tax, stack) · **live-measured pipeline costs from the first real reference-case runs (Aug 30, 2026)** · **Sep 1, 2026 actuals: production provisioning, automated QA decision, attorney-signed eval ledgers, instrumented full-case COGS**

---

## 0. Update log — 2026-09-01 (what changed since prepared)

Material facts that have moved since this case was written; sections below carry dated inline edits.

# 12-Month Business Case — Family Case Review (Sep 2026 – Aug 2027)

---

**Status:** For decision · **Owner:** Product/Finance · **Prepared:** 2026-08-30 · **Updated:** 2026-09-01 **Sources:** snotnoselegal\_market\_study\_mvp\_gtm (market frame \$1, pricing \$6, unit economics \$7, GTM \$8) · mvp\_v1\_prd.md · product\_roadmap.md · ../implementation/mvp\_v1\_implementation\_plan.md (build status & timeline) · PO decisions of Aug 2026 (SLA, pricing, brand, tax, stack) · **live-measured pipeline costs from the first real reference-case runs (Aug 30, 2026)** · **Sep 1, 2026 actuals: production provisioning, automated QA decision, attorney-signed eval ledgers, instrumented full-case COGS**

---

## 0. Update log — 2026-09-01 (what changed since prepared)

Material facts that have moved since this case was written; sections below carry dated inline edits.

1. **Manual QA is removed from the operating model.** The PO decision to eliminate the human-review bottleneck shipped as an automated QA gate ( auto\_qa\_hold\_workflow.md ): screens-complete, minimum-findings, drop-ratio, and render re-verification checks auto-deliver passing reports; failures hold with a customer notice and a 24-hour manual-triage SLA; the founder spot-checks a sample rather than touching every report. AUTO\_APPROVE=1 is armed. This deletes the ~\$12.50/report QA cost line (\$5), removes the ~50-case/month solo throughput ceiling (\$8c, \$12d), and converts Lever 1 of \$8b from a plan into shipped software.
2. **Engineering is complete and production is provisioned** — not "four months to launch-ready." As of Sep 1: full loop proven repeatedly on both reference cases end-to-end (purchase → upload → digitize → analyze → auto-QA → delivered PDF); Render production stack live (Ohio: web, api, Postgres w/ PITR, Redis, clamav — the \$5b ~\$88/mo actuals); promo engine, customer-feedback program, retention tiering (S3 lifecycle), Sentry/PostHog wired and verified. Remaining launch work is **business, not build**: DNS cutover, Stripe live-mode webhook, E&O, counsel sign-offs, credential rotation.
3. **Launch gate 1 is met early:** the reviewing attorney signed **both** reference-case eval ledgers (Sep 1). Current harness state: Gary case **green** 5/5 must-finds (64 findings, auto-delivered as the first live AUTO\_APPROVE case); Brian case 2/4 on the latest fresh run (run-variance remediation — sample count 3 + larger output cap — is scoped and pending a go decision, ~\$12 of re-run spend). The Dec-26 "no eval green, no launch" checkpoint stands for the Brian ledger.
4. **COGS is now instrumented per case, not estimated:** Gary case **\$4.66** all-in tech (model \$3.75 + OCR \$0.91, scanned 602-pp record), Brian case **\$9.42** (born-digital transcript-heavy record, all model, 2-sample union). Batch API validated at ~50% model-cost reduction (\$1.75 Gary harness run). \$5 updated.
5. **Trial GTM in place:** admin-managed 100%-off promo codes with atomic cap enforcement and a Stripe-free \$0 fulfillment path, live-verified — EARLYBIRD (25 uses, early adopters) and SNOT26 (5 uses, trial offer). 30 free slots ≈ **~\$200–300 of COGS as the launch marketing/testimonial budget** — under one month of the \$500 paid-channel floor.

6. **Spanish pulled forward:** the v1.1 "Spanish purchase flow" item partially shipped Sep 1 — bilingual landing, auth, buy-flow disclosures (English-governs banner pending counsel), and a plain-language documents guide, with a per-release parity gate ( `i18n_localization.md` ). The Spanish **report Part A** remains v1.1. Adds a one-time ~\$200–400 professional translation review before Spanish ad spend.
7. **New recurring obligations created:** 24-hour QA-hold triage SLA (founder), promo redemption monitoring, and the counsel review queue now includes the rewritten disclosure set ( `2026-09-01.1` , training-use promise removed), TX/FL/CA disclosure pages, and the Spanish legal copy.

## 1. Executive summary

Snot Nose Legal's **Family Case Review** sells a \$299 flat-fee, AI-analyzed, quality-gated review of a Texas criminal court record to the one persona in the post-conviction market no incumbent serves and who actually pays: the inmate's family. The competitive anchor is not software but the ~\$3,000 **attorney file review** and the \$0 pro se path that burns a family's only writ shot. The market study establishes genuine white space (habeas-specific analysis; consumer access) that none of the five incumbent platforms can enter without abandoning their per-professional-seat licensing.

This case models the next 12 months: four months to launch-ready (per the implementation plan) and eight months of operation. *(Updated 2026-09-01: the build finished early — the full loop is proven end-to-end on both reference cases, production is provisioned on Render, and automated QA is live; the Sep–Dec phase is now legal/business gate work, not engineering. See §0.)*

**The headline numbers (base case):** first-year revenue ≈ **\$36k**, peak cumulative cash need ≈ **\$78k**, monthly contribution turns positive around ~**30 cases/month** at current cost structure — i.e., the business does not self-fund inside 12 months at bootstrap staffing; it buys three things instead: a **proven engine on paid real records**, the **reference metrics** (recall, refund rate, CAC by channel) that de-risk every later tier, and the **referral flywheel** into the Justice/Advocate tiers where the durable revenue lives (roadmap v2/v3). The recommendation (§10) is to proceed, with the pre-agreed kill/pivot checkpoints in §8.

## 2. Problem & product

When a Texas felony conviction becomes final, the family faces a brutal information gap: pay ~\$3,000 for an attorney merely to read the file, pay \$15k–\$50k+ for representation on faith, or let the inmate file pro se — where a weak application burns the effective one shot Texas's subsequent-writ bar allows. **The product sells the decision, not the writ:** every finding grounded to page and line, screened for the six claim families (preserved error, IAC, Brady, 11.073 junk science, sentencing/deadlines, jury selection), passed through an automated QA gate, and delivered as Part A (plain English) + Part B (attorney-ready packet).

**Trust is the product.** The engineering enforces it structurally: findings whose quotes aren't verbatim in the record are dropped by a hard filter; citations are hash-anchored and re-verified at every render; every report clears automated quality gates before delivery — failures hold for human review on a stated 24-hour SLA rather than shipping — with founder spot-checks on a sample; deletion, refunds, and disclosures are audited operations, not promises.

## 3. Market

| Measure | Size      | Definition                                                          |
|---------|-----------|---------------------------------------------------------------------|
| TAM     | \$30B     | Global legal tech & legal AI                                        |
| SAM     | \$1.5–2B  | US criminal defense / public defender legal tech                    |
| SOM     | \$50–100M | Texas post-conviction — the nation's highest Art. 11.07 writ volume |

The consumer wedge is deliberately unsized by incumbents: all five platforms license to lawyers. Texas structural facts driving demand: IAC is the dominant claim; most 11.07 applications are filed pro se; the AEDPA one-year clock is miscalculated fatally and computed by no platform; decades-old scans break incumbents' tooling (our Textract pipeline with per-page confidence is built for exactly these — proven this week on a real 602-page reporter's record).

**Competitive floor:** research is free to the Texas bar (vLex), document review is free to nonprofits (Everlaw for Good). The product never positions as either.

#### 4. Pricing & offer (as decided, Aug 2026)

| Offer                         | Price                                      | Notes                                                             |
|-------------------------------|--------------------------------------------|-------------------------------------------------------------------|
| Family Case Review            | <b>\$299 one-time</b>                      | Uniform across trial and plea lanes (PO decision); 5,000-page cap |
| Page overage                  | +\$49 / 2,500 pages                        | Charged in-flow only when the count crosses the line              |
| Re-run with new documents     | \$99                                       | After report delivery                                             |
| Attorney-signed review add-on | ~\$499                                     | v1.1 (from ~May 2027); licensed TX attorney per review            |
| Payment options               | Stripe-native installments (Affirm/Klarna) | Stripe Tax enabled; TX data-processing determination pending      |

10-business-day SLA from records-complete. Later tiers (Advocate \$99/seat, Chambers \$249/seat, Justice \$0-to-cap, Sovereign site license) are out of this case's revenue but are its strategic payoff.

#### 5. Unit economics

Study budget vs. what the pipeline **actually measured** — updated 2026-09-01 with instrumented full-case costs (CostRecord table) from complete fresh runs of both reference cases at the final configuration (6 screens, context pre-pass, 2-sample union):

| Cost line                 | Study budget    | Measured (Sep 1, 2026)                                                                                    | Note                                                                                                            |
|---------------------------|-----------------|-----------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------|
| OCR                       | ~\$5 / 3,000 pp | <b>\$0.91</b> (Gary: 602 pp scanned × Textract \$1.50/1k); <b>\$0</b> (Brian: born-digital)               | scales ~\$7.50 at the 5,000-page cap                                                                            |
| LLM analysis              | ~\$12           | <b>\$3.75–\$9.42/case</b> (Claude Opus 5; range = scanned vs transcript-heavy record)                     | batch API validated at ~50% off these numbers (\$1.75 Gary harness run); economics gate enforces the batch path |
| Embeddings + storage      | ~\$1            | S3 pennies; 90-day lifecycle tiering applied; embeddings paused by design                                 |                                                                                                                 |
| Payment processing        | ~\$9            | 2.9% + \$0.30                                                                                             | + 0.5% Stripe Tax when taxable; <b>\$0 on promo-trial cases</b> (no Stripe on the free path)                    |
| Human QA                  | ~\$12.50        | <b>\$0 — automated QA gates + founder spot-check sample</b> (auto_qa_hold_workflow.md; AUTO_APPROVE live) | holds route to a 24h manual triage; founder time, not cash                                                      |
| Refund/chargeback reserve | ~\$15           | 5% of revenue                                                                                             | consumer stress purchases                                                                                       |
| <b>Total COGS</b>         | <b>~\$54</b>    | <b>~\$29–34 measured</b> (tech \$5–10 + Stripe ~\$9 + reserve ~\$15)                                      | <b>~89% gross margin at \$299</b>                                                                               |

This case's tables still model the conservative **\$54**; the measured trend is upside and the QA line's removal is structural, not situational (see §8b Lever 1, now shipped).

#### 5b. External tools & services — the documented monthly stack (added 2026-08-31)

The hosting/tooling line items are now provisioned and priced, not estimated (canonical register with tiers and upgrade triggers: docs/operations/external\_services.md):

| Category                           | Services                                                                    | \$/month |
|------------------------------------|-----------------------------------------------------------------------------|----------|
| Hosting (Render, decided platform) | web \$7 · api \$25 · Postgres w/ PITR ~\$20 · Redis \$10 · clamav scan \$25 | ~\$87    |

| Category                                 | Services                                                                                       | \$/month          |
|------------------------------------------|------------------------------------------------------------------------------------------------|-------------------|
| Domain                                   | snotnoselegal.com (~\$15/yr)                                                                   | ~\$1              |
| Email / errors / analytics / uptime / CI | Resend · Sentry · PostHog · UptimeRobot · GitHub — all inside free tiers at launch             | \$0               |
| <b>Fixed tools total</b>                 |                                                                                                | <b>~\$88</b>      |
| Usage-based (in per-case COGS, \$5)      | Anthropic ~\$1.75–3.30/case (batch) · Textract ~\$0.90/case · S3 <\$5/mo · Stripe ~2.9%+\$0.30 | scales with cases |

**Impact on the model:** the solo-founder cash model (\$8c) budgeted **\$150/mo** for hosting; the provisioned reality is **~\$88/mo** — a ~\$62/mo favorable variance that also pre-funds the first paid-tier upgrades (Resend \$20, Sentry \$26 at ~100 cases/mo) without touching the budget line. Cash break-even stays ~4 cases/month; the variance is retained as buffer rather than re-plumbed through the tables.

## 6. 12-month plan

| Phase                       | Months        | Contents                                                                                                                                                                                                                                                                                                                                                                                                                                                                   | Source                                   |
|-----------------------------|---------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------|
| Finish build → launch gates | Sep–Dec 2026  | <b>Build complete as of Sep 1</b> (batch, PDF render, analytics, auto-QA, promos, feedback, Spanish P0 all shipped; production provisioned). Remaining gate work is business-side: Brian-ledger eval green, counsel sign-offs (disclosures 2026-09-01.1, state pages, Spanish legal copy), E&O bound, DNS + Stripe live, credential rotation. <b>Launch can pull forward from Jan-27 if the gates clear early — a favorable schedule variance this case does not spend</b> | implementation plan §3/§7; Sep 1 actuals |
| Launch                      | Jan 2027      | TX-only, SEO + prison-family communities + clinic-declined referrals first; paid search as a capped test                                                                                                                                                                                                                                                                                                                                                                   | market study §8.2                        |
| Ramp                        | Feb–Apr 2027  | 90-day target ~10 paid cases/mo; S0 outcome-mix report validates market sizing                                                                                                                                                                                                                                                                                                                                                                                             | analytics plan §5                        |
| v1.1                        | ~Apr–May 2027 | Spanish Part A + purchase flow; attorney-signed add-on; A/V add-on priced                                                                                                                                                                                                                                                                                                                                                                                                  | roadmap v1.1                             |
| Justice-tier seeding        | Jun–Aug 2027  | Consented high-viability referrals to 2–3 clinics — the v2 gate                                                                                                                                                                                                                                                                                                                                                                                                            | roadmap v2                               |

## 7. Financial model — three scenarios

**Assumptions (labeled):** bootstrap staffing (founder unpaid + part-time senior contractor: \$8k/mo build, \$6k/mo operating — the funded-team alternative per the implementation plan runs ~\$30k/mo and is shown in §8); fixed infra + E&O ~\$600/mo (study §7.4); counsel \$4k one-time (Nov, gate 2) + \$500/mo retainer from launch; marketing = max(\$500/mo test budget, blended CAC \$60 on the ~50% of cases from paid channels — organic channels are the durable base); overage attach 10%, re-run 5%; attorney-signed add-on 8% attach from May 2027 (base/upside only) at \$250 attorney cost; volumes per the study's break-even ladder with the 90-day ~10/mo target hit in March (base).

### Conservative scenario

| Month  | Cases | Revenue | COGS  | Gross | Opex     | EBITDA    | Cumulative |
|--------|-------|---------|-------|-------|----------|-----------|------------|
| Sep-26 | 0     | \$0     | \$0   | \$0   | \$8,600  | \$-8,600  | \$-8,600   |
| Oct-26 | 0     | \$0     | \$0   | \$0   | \$8,600  | \$-8,600  | \$-17,200  |
| Nov-26 | 0     | \$0     | \$0   | \$0   | \$12,600 | \$-12,600 | \$-29,800  |
| Dec-26 | 0     | \$0     | \$0   | \$0   | \$8,600  | \$-8,600  | \$-38,400  |
| Jan-27 | 2     | \$598   | \$108 | \$490 | \$7,600  | \$-7,110  | \$-45,510  |
| Feb-27 | 3     | \$897   | \$162 | \$735 | \$7,600  | \$-6,865  | \$-52,375  |

| Month  | Cases | Revenue | COGS  | Gross   | Opex    | EBITDA   | Cumulative |
|--------|-------|---------|-------|---------|---------|----------|------------|
| Mar-27 | 5     | \$1,495 | \$270 | \$1,225 | \$7,600 | \$-6,375 | \$-58,750  |
| Apr-27 | 7     | \$2,142 | \$378 | \$1,764 | \$7,600 | \$-5,836 | \$-64,586  |
| May-27 | 9     | \$2,740 | \$486 | \$2,254 | \$7,600 | \$-5,346 | \$-69,932  |
| Jun-27 | 11    | \$3,437 | \$594 | \$2,843 | \$7,600 | \$-4,757 | \$-74,689  |
| Jul-27 | 13    | \$4,035 | \$702 | \$3,333 | \$7,600 | \$-4,267 | \$-78,956  |
| Aug-27 | 15    | \$4,682 | \$810 | \$3,872 | \$7,600 | \$-3,728 | \$-82,684  |

**Year totals — revenue \$20,026, EBITDA \$-82,684; peak cumulative cash need \$82,684; first EBITDA-positive month: not in window.**

#### Base scenario

| Month  | Cases | Revenue  | COGS    | Gross   | Opex     | EBITDA    | Cumulative |
|--------|-------|----------|---------|---------|----------|-----------|------------|
| Sep-26 | 0     | \$0      | \$0     | \$0     | \$8,600  | \$-8,600  | \$-8,600   |
| Oct-26 | 0     | \$0      | \$0     | \$0     | \$8,600  | \$-8,600  | \$-17,200  |
| Nov-26 | 0     | \$0      | \$0     | \$0     | \$12,600 | \$-12,600 | \$-29,800  |
| Dec-26 | 0     | \$0      | \$0     | \$0     | \$8,600  | \$-8,600  | \$-38,400  |
| Jan-27 | 3     | \$897    | \$162   | \$735   | \$7,600  | \$-6,865  | \$-45,265  |
| Feb-27 | 6     | \$1,843  | \$324   | \$1,519 | \$7,600  | \$-6,081  | \$-51,346  |
| Mar-27 | 10    | \$3,039  | \$540   | \$2,499 | \$7,600  | \$-5,101  | \$-56,447  |
| Apr-27 | 13    | \$4,035  | \$702   | \$3,333 | \$7,600  | \$-4,267  | \$-60,714  |
| May-27 | 17    | \$5,779  | \$1,168 | \$4,611 | \$7,610  | \$-2,999  | \$-63,713  |
| Jun-27 | 21    | \$7,474  | \$1,634 | \$5,840 | \$7,730  | \$-1,890  | \$-65,603  |
| Jul-27 | 25    | \$8,670  | \$1,850 | \$6,820 | \$7,850  | \$-1,030  | \$-66,633  |
| Aug-27 | 30    | \$10,313 | \$2,120 | \$8,193 | \$8,000  | \$193     | \$-66,440  |

**Year totals — revenue \$42,050, EBITDA \$-66,440; peak cumulative cash need \$66,633; first EBITDA-positive month: Aug-27.**

#### Upside scenario

| Month  | Cases | Revenue  | COGS    | Gross    | Opex     | EBITDA    | Cumulative |
|--------|-------|----------|---------|----------|----------|-----------|------------|
| Sep-26 | 0     | \$0      | \$0     | \$0      | \$8,600  | \$-8,600  | \$-8,600   |
| Oct-26 | 0     | \$0      | \$0     | \$0      | \$8,600  | \$-8,600  | \$-17,200  |
| Nov-26 | 0     | \$0      | \$0     | \$0      | \$12,600 | \$-12,600 | \$-29,800  |
| Dec-26 | 0     | \$0      | \$0     | \$0      | \$8,600  | \$-8,600  | \$-38,400  |
| Jan-27 | 5     | \$1,495  | \$270   | \$1,225  | \$7,600  | \$-6,375  | \$-44,775  |
| Feb-27 | 9     | \$2,740  | \$486   | \$2,254  | \$7,600  | \$-5,346  | \$-50,121  |
| Mar-27 | 14    | \$4,334  | \$756   | \$3,578  | \$7,600  | \$-4,022  | \$-54,143  |
| Apr-27 | 19    | \$5,878  | \$1,026 | \$4,852  | \$7,670  | \$-2,818  | \$-56,961  |
| May-27 | 25    | \$8,670  | \$1,850 | \$6,820  | \$7,850  | \$-1,030  | \$-57,991  |
| Jun-27 | 31    | \$10,612 | \$2,174 | \$8,438  | \$8,030  | \$408     | \$-57,583  |
| Jul-27 | 38    | \$13,253 | \$2,802 | \$10,451 | \$8,240  | \$2,211   | \$-55,372  |
| Aug-27 | 45    | \$15,845 | \$3,430 | \$12,415 | \$8,450  | \$3,965   | \$-51,407  |

**Year totals — revenue \$62,827, EBITDA \$-51,407; peak cumulative cash need \$57,991; first EBITDA-positive month: Jun-27.**

## 8. Cash need, staffing sensitivity, and checkpoints

- **Funding requirement (bootstrap):** peak cumulative need  $\approx$  **\$78k (base)** / \$83k (conservative) / \$72k (upside). Round to a **\$100k envelope** for working-capital comfort (chargebacks, Textract/LLM spikes, counsel overruns).
- **Funded-team alternative:** at the implementation plan's full staffing ( $\sim$ \$30k/mo), the same 12 months need  $\approx$  **\$320–350k**. The bootstrap path is viable precisely because the build is largely done; choose funded staffing only if v1.1/v2 acceleration is worth buying.
- **Sensitivity (base case):** CAC doubling to \$120 costs  $\sim$ \$5k/yr (manageable — contribution/case  $\approx$  \$245); COGS at the measured  $\sim$ \$45 adds  $\sim$ \$1.5k/yr; a \$349 price test (+17%) adds  $\sim$ \$8k/yr revenue at unknown conversion cost — run it only through the \$4.3-compliant experiment path. The model is **volume-dominated**: the single number that matters is paid cases/month.
- **Checkpoints (pre-agreed):**
  - **Dec 2026:** eval harness green + attorney sign-off, or launch slips — do not launch on an unproven engine. (*Status Sep 1: attorney signed both ledgers; Gary case green 5/5; Brian case 2/4 pending the scoped run-variance remediation — this checkpoint is more than half met, three months early.*)
  - **Mar 2027 (90 days post-launch):**  $\geq$  8–10 cases/mo and refund rate  $\leq$  5% → continue; materially below → the S0 outcome-mix data decides pivot (plea-lane emphasis, price, or channel) before more spend.
  - **Jun 2027:**  $\geq$  20% referral opt-in → begin Justice-tier seeding (the v2 gate); below → double down on B2C channels first.

### 8b. The path to profitability (BA answer to "how do we make money on this")

The base case loses money for exactly one structural reason: **fixed operating cost (\$7.6k/mo) needs  $\sim$ 31 cases/month to cover at \$245 contribution, and the ramp only reaches 30 by month 12**. Profitability is therefore engineered from three levers, in order of leverage:

**Lever 1 — lean operations (the big one). Shipped as software, 2026-09-01.** The \$6k/mo operating contractor is 79% of post-launch opex, budgeted for a build that is *already done*. Founder-run operations with an on-call contractor ( $\sim$ \$1.5k/mo) cut fixed cost to  $\sim$ \$3.1k/mo → **break-even falls from 31 to 13 cases/month** — inside the study's own 90-day trajectory. The automated QA gate goes further than the lean-rate assumption: passing reports deliver with zero marginal labor (spot-check sample only), and the failure path is a bounded 24h triage queue, so this lever no longer depends on anyone's discipline — it is the default behavior of the system.

**Lever 2 — contribution stacking (no new customers needed).** Measured COGS (~~\$45, not \$54~~) + ~~overage/re-run attach~~ + the v1.1 attorney-signed add-on pulled forward (10% attach  $\times$  \$249 margin) lifts contribution from \$245 to  **$\approx$ \$289/case** → break-even  $\sim$  **11 cases/month**.

**Lever 3 — Advocate-tier early access (recurring revenue, warm leads).** The QA console is already the attorney workspace dogfood, and the referral list produces warm attorney contacts from month one. An early-access Advocate seat at \$99/mo needs no new engineering to pilot: **10 seats  $\approx$  3 cases of monthly contribution, 25 seats  $\approx$  9 — recurring**. This is the roadmap's Phase-3 revenue sampled early, without gating v2.

**Re-modeled (lean ops + stacked contribution, same base volumes):**

| Post-launch month | Cases | EBITDA        | Cumulative |
|-------------------|-------|---------------|------------|
| Jan-27            | 3     | –\$2,338      | –\$40,738  |
| Feb-27            | 6     | –\$1,527      | –\$42,265  |
| Mar-27            | 10    | –\$511        | –\$42,776  |
| <b>Apr-27</b>     | 13    | <b>+\$350</b> | –\$42,426  |
| May-27            | 17    | +\$1,903      | –\$40,523  |
| Jun-27            | 21    | +\$2,799      | –\$37,724  |

| Post-launch month | Cases | EBITDA   | Cumulative |
|-------------------|-------|----------|------------|
| Jul-27            | 25    | +\$3,695 | -\$34,029  |
| Aug-27            | 30    | +\$5,212 | -\$28,817  |

**Monthly EBITDA turns positive in April 2027 (month 4 post-launch)**, peak cash need drops from ~~\$70k to \*\*\$43k\*\*~~, and the exit run-rate (~\$5k/mo EBITDA and climbing, before any Advocate seats) pays back the build burn during year two. Add 25 early-access Advocate seats and break-even arrives a month earlier with ~\$2.5k/mo of recurring on top.

**What this does NOT change:** the volume dependency. Every lever above lowers the bar; only the channels (clinic-declined referrals at zero CAC, SEO, communities) clear it. The March checkpoint in §8 remains the decision that matters.

### 8c. The solo-founder cash model (the operator's actual position)

Sections 7–8b model *staffed* scenarios from the implementation plan's team assumptions. The operating reality is **one founder, home-based, no payroll** — which converts the model from a P&L question into a cash-flow question, and changes the answer dramatically:

- **Fixed cash ≈ \$1,100/month** (hosting budgeted ~\$150 — provisioned actual ~\$88, see §5b; E&O ~\$250, counsel pay-as-you-go ~\$200, \$500 marketing floor). No salaries, no office.
- **Cash contribution ≈ \$278/case** — cash COGS is ~\$31 (OCR + LLM + storage + Stripe + expected refunds), consistent with the Sep-1 instrumented range (§5). (*Updated 2026-09-01: QA is no longer founder time per case — automated gates deliver passing reports; founder time is a spot-check sample plus the 24h hold-triage queue.*)
- **Cash break-even: ~4 cases/month.**

| Phase                | Cash position                                              |
|----------------------|------------------------------------------------------------|
| Build (Sep–Dec 2026) | –\$4,850 total (hosting + \$4k counsel gate + E&O binding) |
| Jan-27 (3 cases)     | –\$296/mo → cumulative –\$5,146                            |
| Feb-27 (6)           | +\$557 → –\$4,589                                          |
| Mar-27 (10)          | +\$1,629 → –\$2,960                                        |
| <b>Apr-27 (13)</b>   | +\$2,532 → <b>–\$428 (effectively recovered)</b>           |
| May-27 (17)          | +\$4,151 → <b>+\$3,723 cumulative positive</b>             |
| Aug-27 (30)          | +\$8,032/mo → +\$23,273 cumulative                         |

**Read:** on the base volume ramp, the entire venture recovers its cash by ~May 2027 and exits the year at an ~\$8k/month cash run-rate — total capital at risk under **\$6k** plus founder time.

**The honest disclosures that keep this rigorous:** (1) founder time is the real investment — with QA automated (2026-09-01) the old ~50–80 case/month QA ceiling is gone, and the solo binding constraint shifts to **support + hold-triage + everything else**, plausibly 100+ cases/month before paid help returns to the model (a good problem; when it does, it re-enters as support/ops labor, not per-report QA, leaving contribution strong); (2) sweat equity is deferred compensation, not free — the staffed scenarios in §7 remain the true economic cost of the business and the basis for any future hiring or investment conversation; (3) volume remains the gating variable — 4 cases/month is a low bar, but it is still a bar the channels must clear.

### 9. Risks (top 5, from SWOT + study §8.4 + risk register)

| Risk                                                   | Mitigation                                                                                                                                                                       |
|--------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| UPL exposure on the consumer tier                      | Information-report framing throughout (built); counsel gate 2; attorney-signed add-on keeps a human path; never route toward pro se filing                                       |
| Volume doesn't materialize (the model's dominant risk) | Channels are cheap and compounding (SEO, communities, clinic referrals); checkpoints cap the downside; contribution ≈ \$245/case means even small volume is not cash-destructive |

| Risk                                                                         | Mitigation                                                                                                                                                                                                                                                                                  |
|------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Chargebacks/refunds above 5%                                                 | W-2 disclosure archive per case (built — the E-6 evidence packet is one query); 10-day SLA honesty; reserve modeled                                                                                                                                                                         |
| Incumbent module entry                                                       | The moat is the white-space spine + consumer trust surface; incumbents' seat-licensing makes B2C structurally awkward for them                                                                                                                                                              |
| Model-quality regression / eval not green                                    | Recall-first harness on attorney-labeled ledgers is a launch <b>gate</b> , not a hope; live runs already shook out real pipeline defects pre-launch                                                                                                                                         |
| Auto-delivery quality escape (added 2026-09-01: no human reads every report) | Layered gates (grounding filter, render re-verification, screens-complete/min-findings/drop-ratio checks) fail <b>closed</b> into a 24h human-triage hold; founder spot-check sample; per-case feedback survey as the customer-side detector; kill switch is one env var ( AUTO_APPROVE=0 ) |

## 10. KPIs (weekly, from the analytics plan)

Landing → check ≥ 25% · check completion ≥ 80% · fit → ack ≥ 60% · ack → paid ≥ 70% · **\$0 outcome mix** (the market-validation panel) · time-to-records-complete · QA rejection < 10% · refund ≤ 5% · COGS/case ≤ \$54 · north star: **completed reviews that reach a lawyer.**

## 11. Recommendation

**Proceed on the bootstrap path with a \$100k working-capital envelope.** The first year is an investment year by design: the P&L loss (~\$78k base) purchases a proven, eval-gated analysis engine on paid real records, channel economics measured against the study's targets, and the pre-triaged referral pipeline that gates the Justice tier — the sequence the study identifies as the only durable route to the \$50–100M Texas SOM. The two disciplines that keep the downside bounded: the launch gates are hard (no eval green, no launch), and the March volume checkpoint decides pivot with data, not hope.

## 12. Expansion model — two additional high-post-conviction states (BA answer to "what if we expand")

### 12a. State selection

Criteria: post-conviction demand (state prison population as proxy), a document-centric vehicle the engine can screen, a hard deadline that makes the product urgent, records accessibility, and UPL climate. Approximate 2025–26 state prison populations: TX ~134k, CA ~93k, FL ~85k, GA ~50k, OH ~45k, PA ~37k, NY ~33k.

| State                       | Pop. (≈) | Vehicle & deadline                                                 | Fit notes                                                                                                                                                                                                                                                                  |
|-----------------------------|----------|--------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Florida (pick #1)</b>    | 85k      | Rule 3.850 motion, <b>hard 2-year deadline</b>                     | Unified statewide rules, strong e-filing/clerk access, huge pro se volume; deadline engine is directly monetizable. Caution: the Florida Bar is the nation's most aggressive UPL enforcer ( <i>Fla. Bar v. TIKD</i> ) — counsel opinion is a launch gate, not a formality. |
| <b>California (pick #2)</b> | 93k      | § 1473 habeas + resentencing wave (§ 1172.6 felony-murder, SB 483) | Largest addressable population after TX; § 1473(e) (false/repudiated expert evidence) is a statutory analog to TX 11.073 — the junk-science screen ports naturally. Costs: county-by-county practice, fragmented/expensive transcripts, dense free legal-aid ecosystem.    |
| Pennsylvania (alternate)    | 37k      | PCRA, <b>1-year deadline</b>                                       | Best product-fit-per-build-dollar (unified statute, hard clock, active CIU culture) but ~40% of CA's population. Named swap-in if CA build cost or records friction proves heavy.                                                                                          |

### 12b. What a state launch actually costs

Per state: vehicle/checklist/deadline-calendar build + screen-prompt adaptation + state disclosure set (founder time, ~6–8 weeks part-time); UPL counsel opinion ~\$4k; eval corpus — one attorney-labeled reference case, records + labeling ~\$2.5k; foreign qualification/registered agent ~\$500. **One-time ≈ \$7k cash per state**, plus ≈ **\$550/month incremental fixed** while active (\$500 marketing floor + E&O/agent increment). The analysis engine itself is state-agnostic by

construction (grounding, QA gate, event spine); what changes is legal framing — and the **per-state eval ledger green is the same hard launch gate TX has**.

### 12c. Sequencing (gated, solo-founder-realistic)

TX launch Jan-27 unchanged → **FL commit only if the Mar-27 TX checkpoint passes** (build Feb–Apr, launch May-27) → **CA commit only if FL's own 90-day look holds** (build Jun–Aug, launch Sep-27). Expansion ramps modeled at ~60% of the TX curve (population ratio × no-local-network discount).

### 12d. Solo cash model, horizon extended to Feb-28

Assumptions as §8c (\$278/case cash contribution, \$1,100/mo base fixed) plus the per-state costs above; when combined volume crosses 50 cases/month the solo QA ceiling binds and hired QA (\$25/case, Lever 1's lean rate) is applied to **all** cases from that month. *(Updated 2026-09-01: with QA automated, the ~50-case trigger and the \$25/case hire are now a conservatism — the real trigger is support/hold-triage load, later and cheaper. The table below is kept unchanged as the downside-safe version; every row from Oct-27 on is therefore understated.)*

| Month         | TX | FL | CA | Total | Net cash                   | Cumulative                   | TX-only cum. |
|---------------|----|----|----|-------|----------------------------|------------------------------|--------------|
| Jan-27        | 3  | –  | –  | 3     | –\$266                     | –\$5,116                     | –\$5,116     |
| Mar-27        | 10 | –  | –  | 10    | +\$1,680                   | –\$2,868                     | –\$2,868     |
| Apr-27        | 13 | –  | –  | 13    | –\$4,486 (FL launch spend) | –\$7,354                     | –\$354       |
| May-27        | 17 | 2  | –  | 19    | +\$3,632                   | –\$3,722                     | +\$3,272     |
| Aug-27        | 30 | 8  | –  | 38    | +\$1,914 (CA launch spend) | +\$10,460                    | +\$21,100    |
| Oct-27        | 34 | 13 | 4  | 51    | +\$10,703 (QA hire binds)  | +\$31,473                    | +\$37,248    |
| <b>Dec-27</b> | 38 | 18 | 9  | 65    | +\$14,245                  | <b>+\$58,192 (crossover)</b> | +\$55,620    |
| Feb-28        | 42 | 21 | 14 | 77    | +\$17,281                  | <b>+\$91,236</b>             | +\$76,216    |

### 12e. Reads

- **Inside the original 12-month window, expansion is cash-negative relative to staying TX-only** (+\$10.5k vs +\$21.1k by Aug-27): \$14k of launch spend plus early-ramp months. Anyone evaluating on the 12-month frame alone should not expand.
- **The crossover lands ~Dec-27** (8 months after FL launch), and by Feb-28 the expanded business is **+\$15k ahead and diverging ~\$2k+/month**, exiting at a ~\$17.3k/month cash run-rate vs ~\$10.6k TX-only — with SOM roughly 2.3× the TX-only frame and three states of Justice-tier seeding ground.
- **Expansion forces the capacity decision early**: combined volume hits the ~50-case solo ceiling in Oct-27 — a good problem, already priced into the table via hired QA. The founder's binding constraint shifts from demand to throughput a full year earlier than TX-only.
- **The dominant risk is unchanged and now tripled**: volume per state. Each state's Mar-27-style 90-day checkpoint caps its downside at ~\$7k + ~\$550/mo, individually killable without touching the others.

### 12f. Recommendation

**Plan FL now, commit at the Mar-27 checkpoint; hold CA until FL's 90-day read; keep PA as the named alternate.**

The engine and spine are state-agnostic; each expansion buys ~0.6× a Texas at ~\$7k entry cost, with per-state kill switches. Do not let expansion jump the TX launch gates — a second state doubles distribution, not proof.