

Snot Nose Legal

Market Study, MVP Business Case & Go-to-Market Plan

snotnoselegal.com · Texas post-trial appeals & Article 11.07 habeas corpus · August 2026 · Confidential work product. Figures marked [estimated] or [reported] are third-party, not vendor-published.

The verdict in one paragraph. Nobody owns this niche. The research duopoly (Westlaw + CoCounsel, Lexis+ AI) is priced out of reach of the defense bar that does habeas work; the Texas defense bar's research baseline is *free* (vLex Fastcase via State Bar dues); and the innocence/indigent segment can already get Everlaw's ediscovery for \$0 through Everlaw for Good. That kills two obvious positioning moves — competing on research, or competing on generic document review. What no platform offers is the habeas-specific spine: 11.07/AEDPA deadline-and-tolling computation, IAC affidavit tracking, viability triage of inmate intake, a mitigation-record knowledge graph, and cite-to-record drafting for the CCA's mandatory forms. That is Snot Nose Legal's entire product thesis, and it is genuine white space. The pricing anchors are Casefleet (\$30–\$140/seat/mo) below and quote-only enterprise above; the distribution playbook to steal is the bar-benefit / association bundle. The MVP launches on the one persona nobody serves and who actually pays: the inmate's family.

1. Market Frame

Measure	Size	Definition
TAM	\$30B	Global legal tech & legal AI
SAM	\$1.5–2B	US criminal defense / public defender legal tech
SOM	\$50–100M	Texas post-conviction — the highest Article 11.07 writ volume in the nation

Texas structural facts that shape every buyer: 11.07 applications are filed in the convicting court, which has ~180 days before the record moves to the Court of Criminal Appeals; the State answers within 15 days; ineffective-assistance-of-counsel (IAC) is the dominant claim, and most CCA remands happen because trial-counsel affidavits are missing or wrong (TDCAA). Direct

appeal is locked to the four corners of the record; 11.07 requires *building a new evidentiary record* — so tooling must handle both transcript interrogation and investigation artifacts. Post-Michael Morton Act, the Brady disclosure duty extends indefinitely, meaning post-conviction teams re-review entire State files for undisclosed material. Federal habeas adds the 1-year AEDPA clock, tolled by a properly filed state writ — deadline math no platform computes.

2. Feature Matrix

Key: ● shipping · ① partial · ◆ Snot Nose Legal planned per spec · — absent. Planned items come from the spec docs; the codebase is at MVP stage.

2.1 Contested capabilities (incumbents have partial or full coverage)

Capability	Snot Nose Legal	Incumbent coverage
Mass record ingestion & OCR	◆ core pipeline	Everlaw ● (TB-scale); Casefleet ● (per-page metered); Westlaw+CoCounsel ① (batch uploads fail at scale); Lexis ①; vLex ①
A/V transcription (jail calls, bodycam)	◆ Whisper worker (post-MVP add-on)	Everlaw ●; Casefleet ① (\$0.06/min)
Chronology / timeline building	◆ auto from records	Everlaw ● (Storybuilder); Casefleet ● (core product); Westlaw ①; Vincent ①
Cross-transcript inconsistency detection	◆	Everlaw ● (deposition analyzer); Casefleet ① (search + pin-cites); Westlaw ①
Case-law research & citator	① MCP TX-CCA RAG only	Westlaw ● (KeyCite); Lexis ● (Shepard's); vLex ● (Cert) — do not compete here
Cite-to-record verified drafting	◆ side-by-side + Bluebook agent	Westlaw ① (brief check); Lexis ① (drafting); Casefleet ① (pin-cite annotations)
Free / nonprofit path	① Justice tier planned	vLex ● (free via TX bar); Everlaw ● (Everlaw for Good); Lexis — (negotiated only)

2.2 White space (no incumbent offers any of these)

Capability	Snot Nose Legal
Entity / knowledge graph of the record	◆ Neo4j lineage graph (Everlaw offers clustering only; Casefleet entities without a graph view)
11.07 / AEDPA deadline & tolling computation	◆ MCP procedural expert
Exhaustion tracking (state → federal)	◆
Intake triage / viability scoring	◆ clinic dashboard
Junk-science (Art. 11.073) screening	◆ forensic MCP
Clemency / parole packet building	◆
Consumer / family-facing case review	◆ the MVP (\$6–7)
Local / on-prem deployment	◆ Ollama + local stack

Read of the matrix: every capability in table 2.2 is absent from all five incumbents — that white space is the moat. Table 2.1 shows where incumbents are already strong (Everlaw on ingestion and media, Casefleet on chronology, the duopoly on research); Snot Nose Legal should meet the bar there, not try to win there.

3. The Five Competitor Dossiers

3.1 Westlaw Precision + CoCounsel (Thomson Reuters) — \$639/user/mo verified; quote-only above

Features: transcript/record summarization to spot appellate issues; timeline building from witness statements; brief drafting grounded in uploads with clickable Westlaw verification; citation checking; KeyCite citator; agentic "Deep Research." No knowledge graph, no deadline tracking, no media transcription.

Personas: BigLaw and mid-size litigation firms, prosecutors' offices (explicitly marketed), AG offices, appellate boutiques. Thin among solo/indigent defense — cost excludes them.

JTBD: "Find every controlling authority and confirm it's still good law; digest a 5,000-page record fast; check my brief before filing."

Pain solved: research completeness anxiety, cite-check drudgery, transcript-review hours.

Pain caused: quote-only pricing rated "\$\$\$\$\$" on G2; batch uploads of thousands of PDFs fail; Stanford RegLab measured 17–33% hallucination across TR/Lexis AI research tools.

Pricing/licensing: Westlaw Advantage + CoCounsel Essentials \$639/user/mo (verified June 2026); Precision bundles [reported] \$428–\$700/user/mo, stacking toward ~\$3k/seat/mo fully integrated. Per-seat annual contracts, auto-renew, 12–18% multi-year discounts, sales-negotiated above 10 attorneys. No published nonprofit program.

3.2 Lexis+ AI / Protégé (LexisNexis) — [reported] \$128–\$494/user/mo

Features: conversational research with cited answers; upload/summarize pleadings, briefs, matter collections; drafting; real-time Shepard's validation — the strongest citator trust in the market. Gaps: no trial-prep workflow, no chronology builder, no record-analysis pipeline.

Personas: firms of all sizes; research-intensive appellate practices; state offices; prosecution and defense alike.

JTBD: "Authoritative case-law research I can stake a filing on; fast matter-document summaries; drafted memo sections."

Pain solved: citation trust (Shepard's), research speed. **Pain caused:** still ~17% hallucination (best of incumbents, not zero); opaque pricing; AI sold as an add-on outside every base plan.

Pricing/licensing: Protégé unpublished; [reported] \$128–\$494/user/mo all-in, adding ~\$15k–\$59k/yr for a 10-attorney firm over base. Per-seat annual/multi-year; government/education pricing negotiated.

3.3 vLex Fastcase + Vincent AI (Clio) — \$0 via State Bar of Texas dues; Vincent quote-only

Features: free tier: nationwide caselaw/statutes, Texas secondary sources, Cert citator, alerts, Word/Chrome integrations, limited Vincent AI. Paid Vincent: cited research Q&A, argument building, pleading analysis, document comparison, automated timeline extraction, judge/lawyer litigation intelligence, custom workflows.

Personas: solo/small-firm lawyers, court-appointed counsel, legal aid, public defenders — exactly Snot Nose Legal's professional buyer, and every licensed Texas attorney already has it (benefit valued at ~\$1,140/yr).

JTBD: "Do real legal research without a subscription bill; confirm my case is still good law."

Pain solved: the cost of research, entirely. **Pain caused:** AI gated behind an unpriced upgrade; thinner editorial/secondary-source depth than Westlaw/Lexis.

Pricing/licensing: the model that matters — a statewide site license paid by the State Bar of Texas since 2014, upgraded to vLex Fastcase 2024–25 with free member training. This is both the competitive floor ("research is free here") and the distribution playbook worth copying.

3.4 Everlaw + Everlaw for Good — [reported] ~\$2–5k/mo + \$18–35/GB; \$0 for nonprofits

Features: TB-scale processing/OCR; audio/video transcription and translation (jail calls, bodycam, interrogation video); Storybuilder evidence-linked chronologies; deposition transcript annotation and cross-transcript search; EverlawAI deposition analyzer, testimony inconsistency detection. No case-law research, citator, or brief drafting.

Personas: litigation firms and government; via Everlaw for Good: 501(c)(3)s, legal aid, law clinics, CJA panel attorneys, innocence projects (Pennsylvania Innocence Project is a named participant).

JTBD: "Ingest and search an enormous production; find the Brady needle; build the chronology; review AV media without manual transcription."

Pain solved: unstructured terabyte discovery, media transcription, collaboration. **Pain caused:** per-GB hosting "incredibly expensive" — it punishes exactly the decades-long records habeas cases carry; steep learning curve; organizes documents rather than extracting legal narratives.

Pricing/licensing: consumption-based (per-GB), unlimited users, enterprise contracts. Everlaw for Good: free up to a monthly cap; \$6M+ donated across 235 orgs. Hard price anchor: SNL's nonprofit segment can often get Everlaw for \$0.

3.5 Casefleet — \$30–\$140/user/mo (published)

Features: evidence-linked chronologies and visual timelines; entity/witness/issue management; full-text search; transcript viewer with highlight → auto pin-cite; A/V support; OCR. AI tier: document summarization, semantic indexing, named-entity recognition, "Casey"

assistant. No research, citator, or drafting.

Personas: solo and small litigation firms (plaintiff-heavy, marketed to criminal defense), investigators, teams replacing TR Case Notebook. The closest direct product analog.

JTBD: "Turn a mountain of discovery into a defensible chronology; link every fact to its source page; find the contradiction across transcripts."

Pain solved: manual chronology spreadsheets, lost record cites. **Pain caused:** timeline weak until the chronology is manually complete; highlighting degrades on low-quality scans — a critical failure for aged habeas records; "price point somewhat high"; no Bates numbering.

Pricing/licensing: Starter \$30/user/mo (\$23.40 annual, 20 docs/case cap); Advanced AI \$140/user/mo (\$109.20 annual, unlimited docs, 12k doc-intelligence pages + 100k AI credits/user/mo); Enterprise custom at 20+ seats. Overages: OCR \$1/100 pages, transcription \$0.06/min, doc intelligence \$0.01/page. Per-seat SaaS, 22% annual discount, 14-day trial. No nonprofit program.

Runner-ups considered and excluded: TCDLA (association, not software — ~\$100/yr dues, motions/brief banks, CCA argument-prep program; the best partnership channel in the state); Casetext (absorbed into CoCounsel); DISCO/Reduct/Rev (weaker nonprofit story or point solutions); CaseMap (legacy; Casefleet is its replacement); Clearbrief/Midpage/Alexi (no meaningful TX criminal-appellate footprint; Clearbrief's cite-to-record Word add-in is the most interesting adjacent threat); Texas Defender Service / TexasLawHelp / TDCAA (advocacy, pro se help, prosecutor association — context, not competition).

4. Personas & Jobs to Be Done

The spec defines three personas; this study adds a fourth — the inmate and their family, the actual payer in retained work. Only one of the four is served by anyone.

4.1 The Lead Appellate Attorney — partially served (the contested persona)

"When I am drafting an Article 11.07 writ, I want to pull exact, verified citations from the original transcript without breaking my drafting flow, so the court trusts my constitutional claims."

Westlaw/Lexis serve their research needs (at a price), Casefleet their pin-cite needs, Clearbrief their drafting verification. Nobody serves them end-to-end from record to CCA master sheet, and nobody computes their 11.07/AEDPA deadlines or exhaustion path.

4.2 The Mitigation Specialist / Post-Conviction Investigator — unserved

"When I receive a massive dump of institutional records, I want every event mapped into a visual chronological timeline, so I can uncover hidden mitigation factors without reading every page linearly."

Harvey, CoCounsel, and Lexis ignore this persona entirely; Everlaw and Casefleet offer generic chronology with no mitigation semantics (trauma indicators, institutional achievements, TDCJ time credits). The wedge persona into defense teams; the knowledge-graph view is the feature no incumbent has.

4.3 The Clinic Director / Managing Public Defender — unserved

"When we ingest a new batch of 50 inmate requests and raw dockets, I want a summary of legal viability, so I can allocate limited attorney and investigator hours to the cases with the highest probability of success."

No platform triages inmate correspondence or scores 11.07/11.073 viability. This persona also controls the budget at innocence projects and PD offices — the buyer and the white space coincide.

4.4 The Inmate & Their Family — unserved, and the actual payer (MVP persona)

"When my son's direct appeal is denied, I want to know whether a writ has any realistic chance before I spend \$3,000 just for a lawyer to read the file, so I don't drain the family's savings on false hope."

In retained post-conviction work, the check almost always comes from family. Their decision problem is pure cost-versus-probable-outcome, and today it is priced brutally:

- **~\$3,000** — typical flat fee just for an attorney's review of the court documents to assess whether a writ is viable (no filing, no representation).
- **~\$15,000–\$50,000+** [estimated; varies by county and case] — full retained 11.07 representation if the review says "go."
- **\$0 and alone** — the alternative: the majority of 11.07 applications are filed pro se from TDCJ units; weak applications burn the effective one shot Texas's subsequent-writ bar gives them.

The ROI math is the product: a family paying \$3k for a review buys one attorney's afternoon with an unindexed record. An automated case review delivers the same go/no-go signal at a

tenth of the price, and its output doubles as an attorney-ready packet that makes the subsequent \$3k review faster or unnecessary. Every incumbent is structurally unable to serve this persona: they license to lawyers, price per professional seat, and their UX assumes legal training. Funnel effect: a family-purchased review that scores high viability can be forwarded (with consent) to Justice-tier clinics — arriving as a pre-triaged packet instead of a handwritten letter, converting the B2C product into the intake pipeline for the Clinic Director persona.

5. What Texas Practitioners Actually Fight

- **The affidavit bottleneck.** Chronic 11.07 delays waiting on trial-counsel affidavits; attorneys answer without reviewing their files, producing inaccurate affidavits and CCA remands. Affidavit tracking is an unbuilt feature with documented demand on both sides of the bar.
- **Two evidentiary worlds.** Direct appeal = four corners of the record; 11.07 = new extra-record evidence. Every incumbent picks one world; the spec's Escalation Ladder (direct appeal → 11.07 → federal → clemency) is the only design that spans both.
- **Record assembly friction.** Multi-volume reporter's records chased from clerks and court reporters; decades-old scans break competitors' highlighting — OCR quality on aged records is a make-or-break engineering requirement.
- **Brady re-review at scale.** The Michael Morton Act makes the disclosure duty indefinite; post-conviction teams re-audit entire State files. The planned disclosure-timeline "Brady Flag" audit maps directly onto this.
- **The AEDPA clock.** One year from finality, tolled by a properly filed 11.07, equitable tolling rare, miscalculation fatal. A deadline/tolling calculator is unserved across all five platforms — even trial judges get CLE toolkits for managing 11.07 dockets, signaling a B2G opening.

6. Pricing & Licensing Landscape — and Where SNL Lands

Platform	Price point	Licensing model	Nonprofit / bar path
Westlaw + CoCounsel	\$639/user/mo verified; up to ~\$3k	Per-seat annual, auto-renew, quote-driven	None published
Lexis+ AI /	\$128–\$494/user/mo	Per-seat annual, AI as add-	Negotiated only

Protégé	[reported]	on	
vLex Fastcase / Vincent	\$0 (bar) · Vincent quote-only	Statewide bar site license + per-seat upsell	Free to all TX bar members
Everlaw	~\$2–5k/mo + \$18–35/GB [reported]	Consumption (per-GB), unlimited users	Everlaw for Good — free to cap
Casefleet	\$30 / \$140/user/mo published	Per-seat SaaS + usage overages	None

Proposed structure: per-seat, role-aware SaaS with flat case-based ingestion — never per-GB, which punishes decades-long habeas records. Five tiers:

- **Family Case Review (MVP)** — one-time flat ~\$299 per case (no subscription, no per-page metering). Viability report + attorney-ready packet, framed against the ~\$3,000 attorney review anchor: "know before you retain, at a tenth of the cost." Optional ~\$99 re-run when new evidence surfaces. High-viability results route (with consent) to Justice-tier clinics or a vetted-attorney referral list.
- **Advocate** — solo/court-appointed counsel, ~\$99/user/mo (annual ~\$79): above Casefleet Starter, far below the duopoly, buys the habeas spine no one else sells.
- **Chambers** — firms/PD offices, ~\$249/user/mo: full AI analysis, graph, drafting export, Clio/MyCase integrations, pooled AI credits. INVESTIGATOR and VIEWER seats discounted or free — the mitigation-specialist wedge should never be blocked by a paywall.
- **Justice tier** — verified 501(c)(3) innocence projects and clinics, \$0 to a monthly ingestion cap, mirroring Everlaw for Good. Table stakes, not charity: the segment already gets Everlaw free, and these are the reference customers.
- **Sovereign / B2G** — annual site license for self-hosted deployment (the only on-prem option in the market): indigent-defense commissions, county PD offices, courts managing 11.07 dockets. Quote-based; TCDLA/State Bar association bundle as the distribution play.

Licensing terms: monthly/annual with ~20% annual discount (market norm); metered overages only for genuinely marginal costs (A/V minutes, OCR pages at accepted market rates); no auto-renew lock-in beyond one year (a stated grievance against TR); AI included in-tier rather than sold as an add-on (a stated grievance against Lexis).

7. MVP Business Case — Family Tier First

Decision: launch with the Inmate & Family persona only. It has no competitor, the shortest sales cycle (a credit card, not a procurement process), and the product it needs — record in, viability report out — is a straight line through the pipeline already scaffolded (upload → OCR → chunk → entity extraction → analysis workers → report). Professional tiers wait until the consumer product has proven the analysis engine on real records.

7.1 MVP scope

- **Ships:** document upload (PDF/scans), OCR tuned for aged records, three-tier model routing (below), the five viability screens (preserved-error scan, IAC flags, 11.073 junk-science hits, sentencing/time-credit audit, 11.07/AEDPA deadline posture), one deliverable — a plain-English viability report + attorney-ready packet as PDF — plus Stripe checkout and email delivery.
- **Waits:** A/V transcription (the biggest COGS risk; later a priced add-on), the interactive knowledge-graph UI (the graph runs internally; visualization is a professional-tier feature), Clio/MyCase integrations, multi-seat firm workflows, self-hosted deployments.

7.2 AI cost architecture: three-tier model routing

Principle: **cheap models transform, capable models filter** — never place a small model where its misses are unrecoverable. (Full spec: docs/architecture/cost_optimization_ollama.md.)

Tier	Task	Model / delivery	~Cost / 3,000-page case
1 — Mechanical	OCR cleanup, chunk typing, entity candidates	Haiku-class or Gemini Flash, Batch API (−50%)	~\$1
2 — Specialist screens	IAC, Brady, 11.073, sentencing, deadlines (5 passes)	Sonnet-class, Batch API + prompt caching (~0.1× cached reads)	~\$3–4
3 — Synthesis + adjudication	Final synthesis from distilled findings; second engine (Gemini) adversarially verifies findings	Opus-class synthesis; cross-model check at ~5% of dual-pass cost	~\$5–7

only (~50K tokens)

Total ~\$9–12/case vs ~\$40+ for dual full-record passes. Cross-engine disagreement (observed between Gemini and Claude on the two reference case files) becomes a consensus flag routed to human QA; model choices are settled by a recall-first eval harness against attorney-labeled findings ledgers on the reference cases (docs/architecture/model_evaluation.md). Ollama stays in the router abstraction but out of launch spend — a ~\$350/mo GPU only beats ~\$12/case cloud spend at ≥ 150 cases/mo.

7.3 Unit economics per case (\$299 flat fee, ~3,000-page record)

Cost line	Per case	Basis
OCR (3,000 pages)	~\$5	Cloud OCR ~\$1.50/1k pages; self-hosted Tesseract fallback near \$0 [estimated]
LLM analysis	~\$12	Three-tier routing above [estimated]
Embeddings + storage	~\$1	pgvector embeddings $\approx$ \$0.30; S3 $\approx$ \$0.40/case-year at 1.5 GB
Payment processing	~\$9	2.9% + \$0.30 on \$299
Support / human QA	~\$12.50	~30 min/case at \$25/hr loaded — a spot-check on every consumer report is a UPL and quality control, not optional
Refund & chargeback reserve	~\$15	5% of revenue — consumer purchases under emotional stress carry real chargeback risk
Total COGS	~\$54	$\approx$ 82% gross margin at \$299

7.4 Fixed monthly costs at launch (cloud-first)

Line	Monthly	Notes
App hosting (web + API + workers)	\$50–100	Single small VM or PaaS; BullMQ workers colocated
Managed Postgres + pgvector	\$25–70	Neon/RDS small instance

Redis + Neo4j	\$10–75	Upstash Redis ~\$10; Neo4j self-hosted on the app VM at launch (\$0) or Aura ~\$65
S3, email (Resend), Sentry, domain	\$40–75	Mostly free tiers at launch volume
E&O / professional liability insurance	\$150–300	Non-negotiable for a consumer legal-information product
Total fixed	~\$275–620	Plan on ~\$500/mo. One-time pre-launch: attorney review of report templates and disclaimers, ~\$2–5k

7.5 Break-even and cost guardrails

Milestone	Meaning
2 cases/mo	Covers ~\$500 fixed at ~\$245 contribution/case (before marketing)
~6 cases/mo	Covers fixed + a \$500/mo marketing test budget
10 cases/mo	≈ \$2,990 revenue, ≈ \$2,000 net of COGS + fixed
~150 cases/mo	Switch bulk inference to a dedicated GPU (~\$350/mo); below that, cloud APIs beat fixed GPU rent

Three guardrails ensure \$299 always clears cost: (1) a **5,000-page cap** (+\$49 per additional 2,500 pages) converts outlier records from margin erosion into revenue — families can count pages, not gigabytes; (2) **no A/V in the base product** — transcription is the one line that can silently exceed the entire fee; (3) **cloud-first inference** with the Ollama switch as a config change when volume justifies it.

8. Go-to-Market Plan

8.1 Segment sequencing

Phase	Segment	Offer	Gate to next phase
1 — Launch	Inmate families (B2C)	\$299 Family Case Review	Analysis engine proven on paid real records; recall metrics stable on the eval

			harness
2	Innocence projects & clinics	Justice tier (\$0 to cap) seeded by consumer-review referrals	2–3 named reference clinics; triage workflow validated
3	Solo/court-appointed counsel & firms	Advocate \$99 / Chambers \$249 per seat	Referral-list attorneys converting to subscribers
4	B2G / associations	Sovereign self-hosted site license; TCDLA/State Bar bundle	Reference clinics + published outcomes to survive public procurement

8.2 Channels (in order of durability)

- **SEO on the questions families actually search** — "11.07 writ cost," "how to appeal a conviction in Texas," "writ of habeas corpus Texas chances." Cheap, slow, compounding; content is a natural byproduct of the domain expertise in the product.
- **Prison-family communities** — Facebook groups, reentry organizations, TDCJ visitation networks. High trust, zero CPC; requires genuinely useful free content, not ads.
- **Clinic referrals of declined cases** — clinics decline hundreds of requests they cannot staff; a referral to a \$299 review is a service to the family, and high-viability results flow back to the clinic pre-triaged.
- **Vetted-attorney referral list** — monetizable channel that converts into Advocate-tier subscriptions in Phase 3.
- **Paid search (test budget only)** — legal-adjacent CPCs run \$5–15+; a paid-only funnel could push CAC past \$100/case (still profitable at ~\$245 contribution, but not the durable base).
- **Association bundles (Phase 4)** — the vLex/State Bar of Texas playbook: a statewide or TCDLA member benefit is the proven distribution model in this exact market.

8.3 Positioning

- **To families:** "Know before you retain — a complete review of the court record for a tenth of what a lawyer charges to read it."
- **To clinics:** "Your intake backlog, pre-triaged."

- **To counsel:** "The habeas spine no platform sells: deadlines, exhaustion, IAC, Brady, 11.073 — grounded to the record, page and line."
- **Never sound like** "research" (free via the bar) or "document review" (free via Everlaw for Good).

8.4 Risks and honest weaknesses

- **UPL exposure on the consumer tier.** Selling a "probability of relief" opinion to a non-lawyer edges toward unauthorized practice of law. Mitigations: frame the deliverable as an information report, never a recommendation to file; persistent not-legal-advice labeling; route every high-viability result toward a licensed attorney or clinic rather than pro se filing — pushing a weak pro se writ would burn the applicant's effective one shot under the subsequent-writ bar and is the single worst harm the product could cause; optional attorney-signed review add-on (~\$499) keeps a human in the loop.
- **Incumbent module risk.** TR or Lexis could ship a "post-conviction skill," but their pricing keeps them out of this buyer's reach and their pipelines choke on messy bulk uploads. Clearbrief's cite-to-record tech is the likeliest feature to be copied.
- **The free floor.** Any pitch competes with free research (vLex) and free ediscovery (Everlaw for Good); value must live in the habeas-specific spine.
- **Execution gap.** The specs describe a category-defining product; the codebase is MVP with mocked services and an unwritten billing spec. The moat exists only once the planned features ship.
- **Accuracy stakes.** Incumbents hallucinate at 17–33% with billion-dollar databases; in habeas work an error costs a person's liberty. Grounded citations and human-in-the-loop review must be engineering truth before launch, validated by the recall-first eval harness.
- **Brand.** "Snot Nose Legal" is memorable and irreverent; the buyers include grieving families making a \$3k decision and public agencies procuring under scrutiny. A deliberate decision is needed on whether the irreverence sells or screens out the consumer and B2G tiers.

Appendix: Sources

Internal: docs/specifications/ (market_research, use_cases_texas, best_practices, mcp_integration_strategy, auth specs), docs/business_case/ (market_analysis, swot_analysis), docs/design/ (user_journeys, criminal_lifecycle_strategy), docs/architecture/ (expanded_scope_incarceration_reduction, cost_optimization_ollama,

model_evaluation). External: costbench.com (CoCounsel pricing), vaquill.ai (Westlaw/Lexis pricing), arXiv 2405.20362 (Stanford RegLab hallucination study), G2/Capterra reviews, Lawyerist (Lexis/Vincent reviews), Texas Bar Blog (vLex Fastcase member benefit), everlaw.com/everlaw-for-good, casefleet.com/pricing, TDCAA journal (Article 11.07 habeas litigation), Leppard Law (AEDPA tolling), St. Mary's Law Journal (Michael Morton Act). Reported (non-vendor) figures are marked; quote-only vendors do not publish list prices.